

Strongly Polynomial Algorithms for Some Problems Related to Parametric Global Minimum Cuts

Hassene Aissi¹ S. Thomas McCormick² Maurice Queyranne²

¹PSL, Université Paris-Dauphine, LAMSADE

²University of British Columbia, Sauder School of Business

Table of contents

- 1 Parametric
 - The Parametric Problems
- 2 Faster Algorithms for P_{NB}
 - Deterministic
 - Randomized
- 3 Faster Algorithms for P_{max}
 - Deterministic
 - Randomized
- 4 Conclusion

Global Min Cut

- We are given an undirected graph $G = (V, E)$ with non-negative distances (costs) $c_e \in \mathbb{R}^E$.
 - Set $m = |E|$, $n = |V|$ as usual.

Global Min Cut

- We are given an undirected graph $G = (V, E)$ with non-negative distances (costs) $c_e \in \mathbb{R}^E$.
 - Set $m = |E|$, $n = |V|$ as usual.
- A (global) cut is induced by node subset $\emptyset \subset C \subset V$, and includes edges $\delta(C) = \{e \in E \mid \text{exactly one end of } e \text{ is in } C\}$.

Global Min Cut

- We are given an undirected graph $G = (V, E)$ with non-negative distances (costs) $c_e \in \mathbb{R}^E$.
 - Set $m = |E|$, $n = |V|$ as usual.
- A (global) cut is induced by node subset $\emptyset \subset C \subset V$, and includes edges $\delta(C) = \{e \in E \mid \text{exactly one end of } e \text{ is in } C\}$.
- Then a global min cut C^* satisfies $c(\delta(C^*)) = \min_{\text{cuts } C} c(\delta(C))$.

Global Min Cut

- We are given an undirected graph $G = (V, E)$ with non-negative distances (costs) $c_e \in \mathbb{R}^E$.
 - Set $m = |E|$, $n = |V|$ as usual.
- A (global) cut is induced by node subset $\emptyset \subset C \subset V$, and includes edges $\delta(C) = \{e \in E \mid \text{exactly one end of } e \text{ is in } C\}$.
- Then a global min cut C^* satisfies $c(\delta(C^*)) = \min_{\text{cuts } C} c(\delta(C))$.
- Can compute a global min cut in $O(mn + n^2 \log n)$ deterministic time (Stoer-Wagner = SW, Nagamochi-Ibaraki = NI), or $\tilde{O}(n^2)$ randomized time (Karger-Stein = KS), or $\tilde{O}(m)$ randomized time (Karger = K).

Global Min Cut

- We are given an undirected graph $G = (V, E)$ with non-negative distances (costs) $c_e \in \mathbb{R}^E$.
 - Set $m = |E|$, $n = |V|$ as usual.
- A (global) cut is induced by node subset $\emptyset \subset C \subset V$, and includes edges $\delta(C) = \{e \in E \mid \text{exactly one end of } e \text{ is in } C\}$.
- Then a global min cut C^* satisfies $c(\delta(C^*)) = \min_{\text{cuts } C} c(\delta(C))$.
- Can compute a global min cut in $O(mn + n^2 \log n)$ deterministic time (Stoer-Wagner = SW, Nagamochi-Ibaraki = NI), or $\tilde{O}(n^2)$ randomized time (Karger-Stein = KS), or $\tilde{O}(m)$ randomized time (Karger = K).
- There are only $O(n^{\lfloor 2\alpha \rfloor})$ α -approximate min cuts; when $\alpha < \frac{4}{3}$ they can all be computed in $O(n^4)$ deterministic time (NI), or $\tilde{O}(n^{\lfloor 2\alpha \rfloor}) = \tilde{O}(n^2)$ randomized time (KS).

Parametric Global Minimum Cuts

- Now suppose that edge costs are linear functions of d parameters $\mu \in \mathbb{R}^d$.

Parametric Global Minimum Cuts

- Now suppose that edge costs are linear functions of d **parameters** $\mu \in \mathbb{R}^d$.
- Thus we have $d + 1$ edge cost functions c^0, \dots, c^d , and the cost of $e \in E$ w.r.t. μ is $c_\mu(e) = c^0(e) + \sum_{i=1}^d \mu_i c^i(e)$.
 - We do *not* assume that all $c^i(e)$ are non-negative.
 - To avoid NP Hardness, we instead assume that μ is restricted to $M = \{\mu \in \mathbb{R}^d \mid c_\mu(e) \geq 0 \ \forall e \in E\}$.

Parametric Global Minimum Cuts

- Now suppose that edge costs are linear functions of d **parameters** $\mu \in \mathbb{R}^d$.
- Thus we have $d + 1$ edge cost functions c^0, \dots, c^d , and the cost of $e \in E$ w.r.t. μ is $c_\mu(e) = c^0(e) + \sum_{i=1}^d \mu_i c^i(e)$.
 - We do *not* assume that all $c^i(e)$ are non-negative.
 - To avoid NP Hardness, we instead assume that μ is restricted to $M = \{\mu \in \mathbb{R}^d \mid c_\mu(e) \geq 0 \ \forall e \in E\}$.
- Why is parametric global min cut interesting?
 - Models “attack-defend” graph problems where a Defender spends a fixed budget on d resources to reinforce edges against an Attacker.
 - Models situations where costs can change due to external variables.
 - It will turn out to further highlight how the small number of α -approximate solutions leads to more efficient algorithms.

The Global Min Cut Value Function

- Define $Z(\mu)$ to be the cost of a global min cut at μ .
 - Since $Z(\mu)$ is the min of many affine functions (one for each cut), it is a piecewise-linear concave function.
 - AMMQ showed that the number of facets of $Z(\mu)$ is $O(m^d n^2 \log^{d-1} n)$ and they can be computed in $O(m^d \lfloor \frac{d-1}{2} \rfloor n^2 \lfloor \frac{d-1}{2} \rfloor \log^{(d-1) \lfloor \frac{d-1}{2} \rfloor + O(1)} n)$ deterministic time, and $O(mn^4 \log n + n^5 \log^2 n)$ when $d = 1$.
 - When all $c^i(e) \geq 0$, Karger improved this to show that the number of facets of $Z(\mu)$ is $O(n^{d+2})$, and they can be computed in $O(n^{2d+2} \log n)$ randomized time.
- Computing all of $Z(\mu)$ is good, but is maybe too much?

Defining the Parametric Problems

- Computing all of $Z(\mu)$ is good, but is maybe too much?
- E.g., for attack-defend the Attacker only wants to solve $\max_{\mu} Z(\mu)$.
- So define P_{\max} to be the problem of computing the max over μ of $Z(\mu)$ (and an associated global min cut).
- In other applications (e.g. sensitivity analysis) we want to solve P_{NB} :
Given $\mu^0 \in \mathbb{R}^d$ and direction $\nu \in \mathbb{R}^d$, find the next *breakpoint* of $Z(\mu)$ along the ray starting at μ^0 in direction ν .
 - P_{NB} is a sort of *ray-shooting* problem.
 - P_{NB} is effectively a 1-parameter problem, to find the next breakpoint w.r.t. costs $\bar{c}^0 + \mu \bar{c}^1(e)$ with single parameter λ .
- We could solve P_{\max} and P_{NB} by computing $Z(\mu)$, but we want to find something faster.

Megiddo's Parametric Framework

- Megiddo, later with Cohen, gave a black-box way to adapt **linear** algorithms for non-parametric problems to solve parametric problems.
 - Here "linear" means that every comparison is between two affine functions of μ and the data.
- We show that SW is linear, so Megiddo+SW gives an $O(n^{2d+3} \log^d n)$ deterministic algorithm for P_{\max} , and $O(n^5 \log d)$ for P_{NB} .
- Tokuyama saw that KS is linear, so Megiddo+KS gives an $O(n^2 \log^{4d+1} n)$ randomized algorithm for P_{\max} , and $O(n^2 \log^5 n)$ for P_{NB} .
- These are a lot faster than the $O(m^{d \lfloor \frac{d-1}{2} \rfloor} n^{2 \lfloor \frac{d-1}{2} \rfloor} \log^{(d-1) \lfloor \frac{d-1}{2} \rfloor + O(1)} n)$ deterministic and $O(n^{2d+2} \log n)$ randomized algorithms for computing all of $Z(\mu)$.
- However, we'd still like to do better than generic Megiddo.

Summary of Where We Are

Problem	Deterministic	Randomized
Non-param GMC	SW $O(mn + n^2 \log n)$	K $\tilde{O}(m)$ (KS $\tilde{O}(n^2)$)
All $\alpha < \frac{4}{3}$ -approx	NI $O(n^4)$	KS $\tilde{O}(n^2)$
Megiddo $d = 1$	SW $O(n^5 \log n)$	KS $O(n^2 \log^5 n)$
Megiddo gen'l d	SW $O(n^{2d+3} \log^d n)$	KS $O(n^2 \log^{4d+1} n)$
$Z(\mu)$ $d = 1$	$O(mn^4 \log n + n^5 \log^2 n)$	$O(n^4 \log n)$ K
$Z(\mu)$ gen'l d	(big) AMMQ	$O(n^{2d+2} \log n)$ K

Summary of running times so far.

Summary of Where We Are

Problem	Deterministic	Randomized
Non-param GMC	SW $O(mn + n^2 \log n)$	K $\tilde{O}(m)$ (KS $\tilde{O}(n^2)$)
All $\alpha < \frac{4}{3}$ -approx	NI $O(n^4)$	KS $\tilde{O}(n^2)$
Megiddo $d = 1$	SW $O(n^5 \log n)$	KS $O(n^2 \log^5 n)$
Megiddo gen'l d	SW $O(n^{2d+3} \log^d n)$	KS $O(n^2 \log^{4d+1} n)$
$Z(\mu)$ $d = 1$	$O(mn^4 \log n + n^5 \log^2 n)$	$O(n^4 \log n)$ K
$Z(\mu)$ gen'l d	(big) AMMQ	$O(n^{2d+2} \log n)$ K

Big gap between non-parametric and computing all of $Z(\mu)$ running times, even for $d = 1$

Summary of Where We Are

Problem	Deterministic	Randomized
Non-param GMC	SW $O(mn + n^2 \log n)$	K $\tilde{O}(m)$ (KS $\tilde{O}(n^2)$)
All $\alpha < \frac{4}{3}$ -approx	NI $O(n^4)$	KS $\tilde{O}(n^2)$
Megiddo $d = 1$	SW $O(n^5 \log n)$	KS $O(n^2 \log^5 n)$
Megiddo gen'l d	SW $O(n^{2d+3} \log^d n)$	KS $O(n^2 \log^{4d+1} n)$
$Z(\mu)$ $d = 1$	$O(mn^4 \log n + n^5 \log^2 n)$	$O(n^4 \log n)$ K
$Z(\mu)$ gen'l d	(big) AMMQ	$O(n^{2d+2} \log n)$ K

Smaller gap between non-parametric and Megiddo running times (compare to $Z(\mu)$ times in blue); for $d = 1$, KS gap is just logs. Note that using Megiddo to solve P_{NB} is just general Megiddo with d set to 1.

Summary of Where We Are

Problem	Deterministic	Randomized
Non-param GMC	SW $O(mn + n^2 \log n)$	K $\tilde{O}(m)$ (KS $\tilde{O}(n^2)$)
All $\alpha < \frac{4}{3}$ -approx	NI $O(n^4)$	KS $\tilde{O}(n^2)$
Megiddo $d = 1$	SW $O(n^5 \log n)$	KS $O(n^2 \log^5 n)$
Megiddo gen'l d	SW $O(n^{2d+3} \log^d n)$	KS $O(n^2 \log^{4d+1} n)$
$Z(\mu)$ $d = 1$	$O(mn^4 \log n + n^5 \log^2 n)$	$O(n^4 \log n)$ K
$Z(\mu)$ gen'l d	(big) AMMQ	$O(n^{2d+2} \log n)$ K
P_{NB} ($\sim d = 1$)	???	???
P_{max} (\sim gen'l d)	???	???

Hoped-for results in this paper in **red**. Compare to non-param lower bounds in **green**, various upper bounds in **blue**.

Outline


- 1 Parametric
 - The Parametric Problems
- 2 Faster Algorithms for P_{NB}
 - Deterministic
 - Randomized
- 3 Faster Algorithms for P_{\max}
 - Deterministic
 - Randomized
- 4 Conclusion

Using Stoer-Wagner to Solve P_{NB}

- SW finds a node ordering v_1, \dots, v_n such that (v_{n-1}, v_n) is a **pendent pair**, i.e., either $\delta(v_n)$ is a global min cut, or we can contract edge $\{v_{n-1}, v_n\}$ without losing any optimal cuts.


Using Stoer-Wagner to Solve P_{NB}

- SW finds a node ordering v_1, \dots, v_n such that (v_{n-1}, v_n) is a **pendent pair**, i.e., either $\delta(v_n)$ is a global min cut, or we can contract edge $\{v_{n-1}, v_n\}$ without losing any optimal cuts.
- Let G^r be contracted graph at iteration r . Define $Z^r(\mu)$ to be min of $\bar{c}(\delta(v))$ for $v \in V^r$ and compute λ^r like:



Using Stoer-Wagner to Solve P_{NB}


- SW finds a node ordering v_1, \dots, v_n such that (v_{n-1}, v_n) is a **pendent pair**, i.e., either $\delta(v_n)$ is a global min cut, or we can contract edge $\{v_{n-1}, v_n\}$ without losing any optimal cuts.
- Let G^r be contracted graph at iteration r . Define $Z^r(\mu)$ to be min of $\bar{c}(\delta(v))$ for $v \in V^r$ and compute λ^r like:



- Update an UB $\bar{\lambda}$ on λ_{NB} by λ_r , and do SW to find and contract a pendent pair w.r.t. $\bar{\lambda}$; since $Z(\lambda)$ is concave, λ^r upper bounds λ_{NB} .

Using Stoer-Wagner to Solve P_{NB}

- SW finds a node ordering v_1, \dots, v_n such that (v_{n-1}, v_n) is a **pendent pair**, i.e., either $\delta(v_n)$ is a global min cut, or we can contract edge $\{v_{n-1}, v_n\}$ without losing any optimal cuts.
- Let G^r be contracted graph at iteration r . Define $Z^r(\mu)$ to be min of $\bar{c}(\delta(v))$ for $v \in V^r$ and compute λ^r like:



- Update an UB $\bar{\lambda}$ on λ_{NB} by λ_r , and do SW to find and contract a pendent pair w.r.t. $\bar{\lambda}$; since $Z(\lambda)$ is concave, λ^r upper bounds λ_{NB} .
- This is correct, and runs in same $O(mn + n^2 \log n)$ time as SW.

Summary of Running Times

Problem	Deterministic	Randomized
Non-param GMC	SW $O(mn + n^2 \log n)$	K $\tilde{O}(m)$ (KS $\tilde{O}(n^2)$)
All $\alpha < \frac{4}{3}$ -approx	NI $O(n^4)$	KS $\tilde{O}(n^2)$
Megiddo $d = 1$	SW $O(n^5 \log n)$	KS $O(n^2 \log^5 n)$
Megiddo gen'l d	SW $O(n^{2d+3} \log^d n)$	KS $O(n^2 \log^{4d+1} n)$
$Z(\mu)$ $d = 1$	$O(mn^4 \log n + n^5 \log^2 n)$	$O(n^4 \log n)$ K
$Z(\mu)$ gen'l d	(big) AMMQ	$O(n^{2d+2} \log n)$ K
P_{NB} ($\sim d = 1$)	SW $O(mn + n^2 \log n)$???
P_{\max} (\sim gen'l d)	???	???

Here we saved a lot w.r.t. Megiddo, and matched the non-parametric lower bound.

Using Karger-Stein to Solve P_{NB}


- KS selects an edge to contract randomly, proportional to its cost.

Using Karger-Stein to Solve P_{NB}

- KS selects an edge to contract randomly, proportional to its cost.
- After contracting to 2 nodes, KS show that the remaining induced cut is a min cut with probability at least $1/\binom{n}{2}$, and this can be put into a framework that will identify a min cut with high probability.


Using Karger-Stein to Solve P_{NB}

- KS selects an edge to contract randomly, proportional to its cost.
- After contracting to 2 nodes, KS show that the remaining induced cut is a min cut with probability at least $1/\binom{n}{2}$, and this can be put into a framework that will identify a min cut with high probability.
- Compute μ^r like this:



Using Karger-Stein to Solve P_{NB}

- KS selects an edge to contract randomly, proportional to its cost.
- After contracting to 2 nodes, KS show that the remaining induced cut is a min cut with probability at least $1/\binom{n}{2}$, and this can be put into a framework that will identify a min cut with high probability.
- Compute μ^r like this:



- Choose e to contract with probability proportional to $c_{\mu^r}(e)$; since $Z(\mu)$ is concave, μ^r upper bounds μ_{NB} .

Using Karger-Stein to Solve P_{NB}

- Compute μ_{NB} as the intersection of the final cut line and $L(\mu)$, and repeat in the KS framework.

Using Karger-Stein to Solve P_{NB}

- Compute μ_{NB} as the intersection of the final cut line and $L(\mu)$, and repeat in the KS framework.
- As in KS, when μ_{NB} exists, the probability that a global min cut survives all the contractions is at least $1/\binom{n}{2}$; if there is no breakpoint in direction ν , then the algorithm recognizes this with probability one.

Using Karger-Stein to Solve P_{NB}

- Compute μ_{NB} as the intersection of the final cut line and $L(\mu)$, and repeat in the KS framework.
- As in KS, when μ_{NB} exists, the probability that a global min cut survives all the contractions is at least $1/\binom{n}{2}$; if there is no breakpoint in direction ν , then the algorithm recognizes this with probability one.
- Thus using the KS framework is correct, and runs in same $\tilde{O}(n^2)$ time as KS.

Using Karger-Stein to Solve P_{NB}

- Compute μ_{NB} as the intersection of the final cut line and $L(\mu)$, and repeat in the KS framework.
- As in KS, when μ_{NB} exists, the probability that a global min cut survives all the contractions is at least $1/\binom{n}{2}$; if there is no breakpoint in direction ν , then the algorithm recognizes this with probability one.
- Thus using the KS framework is correct, and runs in same $\tilde{O}(n^2)$ time as KS.
- There is a minor technical point about how to implement the random edge contractions: Here the parametric costs interfere with the KS matrix update technique, but we can replace the static matrices with separate matrices for \bar{c}^0 and \bar{c}^1 to achieve the same effect.

Summary of Running Times

Problem	Deterministic	Randomized
Non-param GMC	SW $O(mn + n^2 \log n)$	K $\tilde{O}(m)$ (KS $\tilde{O}(n^2)$)
All $\alpha < \frac{4}{3}$ -approx	NI $O(n^4)$	KS $\tilde{O}(n^2)$
Megiddo $d = 1$	SW $O(n^5 \log n)$	KS $O(n^2 \log^5 n)$
Megiddo gen'l d	SW $O(n^{2d+3} \log^d n)$	KS $O(n^2 \log^{4d+1} n)$
$Z(\mu)$ $d = 1$	$O(mn^4 \log n + n^5 \log^2 n)$	$O(n^4 \log n)$ K
$Z(\mu)$ gen'l d	(big) AMMQ	$O(n^{2d+2} \log n)$ K
P_{NB} ($\sim d = 1$)	SW $O(mn + n^2 \log n)$	KS $O(n^2 \log^3 n)$
P_{\max} (\sim gen'l d)	???	???

Here we saved only \log factors w.r.t. Megiddo, but that's all the gap we had to work with; our ideas don't seem to extend to **Karger's improvement**.

Outline

- 1 Parametric
 - The Parametric Problems
- 2 Faster Algorithms for P_{NB}
 - Deterministic
 - Randomized
- 3 Faster Algorithms for P_{\max}
 - Deterministic
 - Randomized
- 4 Conclusion

Solving P_{\max} : Overview and Techniques

- Following Mulmuley and AMMQ we want to use two ideas to compute $\mu^* = \max_{\mu} Z(\mu)$:
 - 1 Approximate duality between global MC and max spanning tree.
 - 2 Ability to compute all $O(n^2)$ α -approximate solutions for $\alpha < \frac{4}{3}$.

Solving P_{\max} : Overview and Techniques

- Following Mulmuley and AMMQ we want to use two ideas to compute $\mu^* = \max_{\mu} Z(\mu)$:
 - ① Approximate duality between global MC and max spanning tree.
 - ② Ability to compute all $O(n^2)$ α -approximate solutions for $\alpha < \frac{4}{3}$.
- But max spanning tree makes sense only when all costs are linearly ordered, and parametric costs typically are not.

Solving P_{\max} : Overview and Techniques

- Following Mulmuley and AMMQ we want to use two ideas to compute $\mu^* = \max_{\mu} Z(\mu)$:
 - ① Approximate duality between global MC and max spanning tree.
 - ② Ability to compute all $O(n^2)$ α -approximate solutions for $\alpha < \frac{4}{3}$.
- But max spanning tree makes sense only when all costs are linearly ordered, and parametric costs typically are not.
- And we need to narrow down our search for μ^* to a region small enough that the α -approximate min cuts include all cuts defining μ^* .

Solving P_{\max} : Overview and Techniques

- Following Mulmuley and AMMQ we want to use two ideas to compute $\mu^* = \max_{\mu} Z(\mu)$:
 - ① Approximate duality between global MC and max spanning tree.
 - ② Ability to compute all $O(n^2)$ α -approximate solutions for $\alpha < \frac{4}{3}$.
- But max spanning tree makes sense only when all costs are linearly ordered, and parametric costs typically are not.
- And we need to narrow down our search for μ^* to a region small enough that the α -approximate min cuts include all cuts defining μ^* .
- We use a technique from computational geometry called **point location in arrangements** (PLA) to achieve both of these.

Solving P_{\max} : Overview and Techniques


- Following Mulmuley and AMMQ we want to use two ideas to compute $\mu^* = \max_{\mu} Z(\mu)$:
 - 1 Approximate duality between global MC and max spanning tree.
 - 2 Ability to compute all $O(n^2)$ α -approximate solutions for $\alpha < \frac{4}{3}$.
- But max spanning tree makes sense only when all costs are linearly ordered, and parametric costs typically are not.
- And we need to narrow down our search for μ^* to a region small enough that the α -approximate min cuts include all cuts defining μ^* .
- We use a technique from computational geometry called **point location in arrangements** (PLA) to achieve both of these.
- In PLA we are given:
 - a set \mathcal{H} of hyperplanes (think the μ s.t. $c_{\mu}(e) = c_{\mu}(e')$);
 - a polytope P (think the region M where all $c_{\mu}(e) \geq 0$); and
 - an unknown target (think μ^*).

Solving P_{\max} : Overview and Techniques

- Following Mulmuley and AMMQ we want to use two ideas to compute $\mu^* = \max_{\mu} Z(\mu)$:
 - 1 Approximate duality between global MC and max spanning tree.
 - 2 Ability to compute all $O(n^2)$ α -approximate solutions for $\alpha < \frac{4}{3}$.
- But max spanning tree makes sense only when all costs are linearly ordered, and parametric costs typically are not.
- And we need to narrow down our search for μ^* to a region small enough that the α -approximate min cuts include all cuts defining μ^* .
- We use a technique from computational geometry called **point location in arrangements** (PLA) to achieve both of these.
- In PLA we are given:
 - a set \mathcal{H} of hyperplanes (think the μ s.t. $c_{\mu}(e) = c_{\mu}(e')$);
 - a polytope P (think the region M where all $c_{\mu}(e) \geq 0$); and
 - an unknown target (think μ^*).
- Then the task is to find a simplex in a cell of $\mathcal{H} \cap P$ containing μ^* .

Weak Duality between GMC and Max Spanning Tree

- Define \mathcal{H}_1 as the set of $O(m^2)$ hyperplanes where $c_\mu(e) = c_\mu(e')$ and run PLA for $(\mathcal{H}_1, M, \mu^*)$ to get simplex S_1 .



- By the definition of \mathcal{H}_1 and PLA, we know that $\mu^* \in S_1$ and all $c_\mu(e)$ are linearly ordered for $\mu \in S_1$.

Weak Duality between GMC and Max Spanning Tree


- By the definition of \mathcal{H}_1 and PLA, we know that $\mu^* \in S_1$ and all $c_\mu(e)$ are linearly ordered for $\mu \in S_1$.
- Thus we can compute a max spanning tree T in S_1 .
- Let \bar{e} be a min-cost edge in T .
 - Since every cut hits T we get $Z(\mu^*) \geq c_\mu(\bar{e})$ for all $\mu \in S_1$.
 - Let \bar{C} be the fundamental cut in $T - \bar{e}$; since T is a MST we have $Z(\mu^*) \leq c_{\mu^*}(\bar{C}) \leq mc_{\mu^*}(\bar{e})$.
 - Thus $c_{\mu^*}(\bar{e}) \leq Z(\mu^*) \leq mc_{\mu^*}(\bar{e})$, and so $c_{\mu^*}(\bar{e})$ is a fairly tight estimate of $Z(\mu^*)$.
- Now we need to use PLA a second time to further narrow in on μ^* so we can get the cuts inducing it via α -approximate cuts.

Narrowing in on α -Approximate Cuts

- Choose


- $\bar{\alpha}$ s.t. $1 < \bar{\alpha} < \sqrt{\frac{4}{3}}$ (note: $0 < \frac{\bar{\alpha}^2 - 1}{m} < 1$);
- $p = 1 + \lceil \log \frac{m^2}{\bar{\alpha}^2 - 1} / \log \bar{\alpha}^2 \rceil$ so that $\frac{\bar{\alpha}^2 - 1}{m} \bar{\alpha}^{2(p-1)} > m$ (note: $p = O(\log n)$);
- $g_i(\bar{e}, \mu) = \frac{\bar{\alpha}^2 - 1}{m} \bar{\alpha}^{2(i-1)} c_\mu(\bar{e})$ for $i = 1, \dots, p$, $g_0(\bar{e}, \mu) = 0$ (note: $g_1(\bar{e}, \mu) < c_\mu(\bar{e})$ and $g_p(\bar{e}, \mu) > m c_\mu(\bar{e})$).

- Define \mathcal{H}_2 as the $O(m \log n)$ hyperplanes where $c_\mu(e) = g_i(\bar{e}, \mu)$, $\forall e \in E$, $i = 1, \dots, p$, and set $S_2 = \text{PLA}(\mathcal{H}_2, S_1, \mu^*)$:



Computing Min Cuts and μ^*

- Due to how we defined the $g_i(\bar{e}, \mu)$, we know that any cut defining μ^* must be an $\bar{\alpha}$ -approximate cut for any $\mu \in S_2$.
- Thus we could compute the $O(n^2)$ $\bar{\alpha}$ -approximate cuts in \mathcal{C} and compute their lower envelope to get μ^* , but this would take $\Omega(n^{2d})$ time, too slow.
- Instead, define \mathcal{H}_3 as the $O(n^4)$ hyperplanes where $c_\mu(C) = c_\mu(C')$ for $C, C' \in \mathcal{C}$ and set $S_3 = \text{PLA}(\mathcal{H}_3, S_2, \mu^*)$.



Computing Min Cuts and μ^*

- Since μ^* is the intersection of d cuts in \mathcal{C} , it must be a vertex of S_3 , and so this last call of PLA finds μ^* more efficiently.
- PLA is a recursive procedure; when we solve the recursion, we get the claimed $O(n^4 \log^{d-1} n)$ running time.
- I skipped a technicality that arises when $c_\mu(\bar{e}) = 0$ for some $\mu \in S_1$.

Summary of Running Times

Problem	Deterministic	Randomized
Non-param GMC	SW $O(mn + n^2 \log n)$	K $\tilde{O}(m)$ (KS $\tilde{O}(n^2)$)
All $\alpha < \frac{4}{3}$ -approx	NI $O(n^4)$	KS $\tilde{O}(n^2)$
Megiddo $d = 1$	SW $O(n^5 \log n)$	KS $O(n^2 \log^5 n)$
Megiddo gen'l d	SW $O(n^{2d+3} \log^d n)$	KS $O(n^2 \log^{4d+1} n)$
$Z(\mu)$ $d = 1$	$O(mn^4 \log n + n^5 \log^2 n)$	$O(n^4 \log n)$ K
$Z(\mu)$ gen'l d	(big) AMMQ	$O(n^{2d+2} \log n)$ K
P_{NB} ($\sim d = 1$)	SW $O(mn + n^2 \log n)$	KS $O(n^2 \log^3 n)$
P_{\max} (\sim gen'l d)	$O(n^4 \log^{d-1} n)$???

We saved a lot compared to Megiddo, but even for $d = 1$ still much slower than our deterministic P_{NB} algorithm, suggesting that P_{\max} for $d = 1$ is strictly harder than P_{NB} .

Summary of Running Times

Problem	Deterministic	Randomized
Non-param GMC	SW $O(mn + n^2 \log n)$	K $\tilde{O}(m)$ (KS $\tilde{O}(n^2)$)
All $\alpha < \frac{4}{3}$ -approx	NI $O(n^4)$	KS $\tilde{O}(n^2)$
Megiddo $d = 1$	SW $O(n^5 \log n)$	KS $O(n^2 \log^5 n)$
Megiddo gen'l d	SW $O(n^{2d+3} \log^d n)$	KS $O(n^2 \log^{4d+1} n)$
$Z(\mu)$ $d = 1$	$O(mn^4 \log n + n^5 \log^2 n)$	$O(n^4 \log n)$ K
$Z(\mu)$ gen'l d	(big) AMMQ	$O(n^{2d+2} \log n)$ K
P_{NB} ($\sim d = 1$)	SW $O(mn + n^2 \log n)$	KS $O(n^2 \log^3 n)$
P_{\max} (\sim gen'l d)	$O(n^4 \log^{d-1} n)$???

Notice that running time for our P_{\max} algorithm is just log factors more than for computing all $\bar{\alpha}$ -approximate min cuts.

Solving P_{\max} Randomly

- We have a faster randomized algorithm than the algorithm of Tokuyama but only for $d = 1$.
- So far we don't know how to do for $d > 1 \dots$

Final Summary of Running Times

Problem	Deterministic	Randomized
Non-param GMC	SW $O(mn + n^2 \log n)$	K $\tilde{O}(m)$ (KS $\tilde{O}(n^2)$)
All $\alpha < \frac{4}{3}$ -approx	NI $O(n^4)$	KS $\tilde{O}(n^2)$
Megiddo $d = 1$	SW $O(n^5 \log n)$	KS $O(n^2 \log^5 n)$
Megiddo gen'l d	SW $O(n^{2d+3} \log^d n)$	KS $O(n^2 \log^{4d+1} n)$
$Z(\mu)$ $d = 1$	$O(mn^4 \log n + n^5 \log^2 n)$	$O(n^4 \log n)$ K
$Z(\mu)$ gen'l d	(big) AMMQ	$O(n^{2d+2} \log n)$ K
P_{NB} ($\sim d = 1$)	SW $O(mn + n^2 \log n)$	KS $O(n^2 \log^3 n)$
P_{\max} (\sim gen'l d)	$O(n^4 \log^{d-1} n)$???

New results in this paper in red. Compare to non-param lower bounds in green, various upper bounds in blue.

Conclusion

- Solving P_{NB} and P_{max} by computing $Z(\mu)$ is slow.

Conclusion

- Solving P_{NB} and P_{max} by computing $Z(\mu)$ is slow.
- We could use Megiddo+SW to solve them faster deterministically, or Megiddo+KS to solve them faster randomly, which give the impression that P_{NB} and P_{max} for $d = 1$ have the same complexity.
 - Our algorithms suggest that P_{NB} is easier than P_{max} for $d = 1$.

Conclusion

- Solving P_{NB} and P_{max} by computing $Z(\mu)$ is slow.
- We could use Megiddo+SW to solve them faster deterministically, or Megiddo+KS to solve them faster randomly, which give the impression that P_{NB} and P_{max} for $d = 1$ have the same complexity.
 - Our algorithms suggest that P_{NB} is easier than P_{max} for $d = 1$.
- We propose specialized algorithms for solving P_{NB} and P_{max} that are significantly faster than Megiddo.
 - The P_{NB} algorithms are essentially as fast as the non-parametric algorithms.
 - The deterministic P_{max} algorithm further elaborates computational geometry techniques and is much faster than Megiddo+SW.

Conclusion

- Solving P_{NB} and P_{max} by computing $Z(\mu)$ is slow.
- We could use Megiddo+SW to solve them faster deterministically, or Megiddo+KS to solve them faster randomly, which give the impression that P_{NB} and P_{max} for $d = 1$ have the same complexity.
 - Our algorithms suggest that P_{NB} is easier than P_{max} for $d = 1$.
- We propose specialized algorithms for solving P_{NB} and P_{max} that are significantly faster than Megiddo.
 - The P_{NB} algorithms are essentially as fast as the non-parametric algorithms.
 - The deterministic P_{max} algorithm further elaborates computational geometry techniques and is much faster than Megiddo+SW.
- Open question:
 - There should be a faster, specialized, randomized algorithm for P_{max} .

Any questions?

Questions?

Comments?